HarperCollins Publishers, LLC
William S. Adams

Comment from HarperCollins Publishers

Comment ID: COLC-2023-0006-10327

Artificial Intelligence and Copyright

Notice of Inquiry and Request for Comments

(Docket No. 2023-6)

HarperCollins Publishers welcomes the opportunity to provide these Reply Comments to the Notice of Inquiry and Request for Comments on the intersection of copyright law and artificial intelligence. HarperCollins is not only the oldest trade publisher in the United States but is also keenly focused on the latest developments in technology and copyright law. Our comments below are fully supportive of and consistent with those contained in the Submission of the American Association of Publishers and reiterate its conclusions:

1. That copying and use of copyrighted material for the training of Gen AI requires consent;
2. That unauthorized use of copyrighted material runs counter to the policy of copyright law;
3. That Gen AI companies must be transparent about the specific works of authorship used by them in training and tuning their models; and
4. That a legislative solution may be essential if courts fail to protect copyright owners from widespread and unauthorized use of their works by Gen AI companies.

Our brief comments below are in opposition to much of the post-hoc self-justification expressed by the Gen AI companies that seek to profit off the expressive content of authors.

Imbedded in the Constitution, copyright law has at its foundation the objective of promoting the development of the useful arts by granting a property right to authors and others who invest in creating them. It balances the interests of the creators with those of later artists and writers, encouraging and protecting the former while allowing the accretion of new works that may in certain circumstances make limited use of the prior work.

Gen AI companies threaten to upend this balance by reading or ingesting virtually all books ever written, then producing a torrent of computer-generated material at previously unimagined speed and volume. Nevertheless, in their submissions, proponents of Gen AI claim that their AI models do not rely on the expressive content of the thousands of authors' works they have used for training and finetuning their models. Their submissions minimize the importance of the expressive content of the books they copy, creatively describing their use in other ways. For some, it is akin to mining: "extract[ing] facts and statistical patterns" (Andreessen Horowitz at 6), or "extracting unprotectable elements" (Anthropic at 7), or "extracting unprotected information *about* the English language – i.e., the correlations, patterns, and relationships among the 26 letters of the alphabet and the 1 million English language words" (BSA at 8).

As if these are minerals lying deep under the surface of some previously uncharted and unclaimed territory.

In other submissions they study the expressive works themselves, "analyz[ing] the structure and syntax of language . . . to discern the statistical relationships between words." (OpenAI at 11). Or they use the expressive texts to "generat[e] [vectors] to represent not just words but information about the semantic and contextual meaning of the words and their relationships to other words in the vocabulary." In other words, they use the underlying expressive texts to generate *new* expressive texts. (Microsoft at 6).

In fact, what all these examples are extracting or analyzing is the way in which authors have expressed themselves – precisely so the Gen AI companies can emulate that expression convincingly for AI users. Rather than create their own sentences, paragraphs and stories to train their systems or limit themselves to using public domain materials, the Gen AI companies have pirated the copyrighted work of others and now strain to recharacterize their cost-cutting short-cuts.

Moreover, the Gen AI companies not only are using authors' books for initial training but will remain reliant on the ongoing use of the underlying expression created by authors. Without the continuous training provided through the use of human-created works, Gen AI-generated writing that trains only on its own output rapidly deteriorates into "gibberish." Rahul Rao, AI-Generated Data Can Poison Future AI Models, *Scientific America*, July 28, 2023. Not only have the Gen AI companies ingested copyrighted works to create their models, but they must continue to rely on the copyrighted works of human authors who have not given their consent to prevent the Gen AI models from corrupting their own output with inferior Gen AI-generated content.

We appreciate the opportunity to provide our views and look forward to working with the Copyright Office on this matter.

HarperCollins Publishers, LLC

December 6, 2023

Comment from LUCITA

Comment ID: COLC-2023-0006-6626

Dear US Copyright Office,

Thank you for allowing the public to comment on the matter of generative AI and copyright. I am a published literary author, journalist and researcher, publisher, and a communications and marketing professional. I also own a small business which produces original content for clients. This means I have worn numerous hats throughout my career, and have had experience with and insight into the creation and usage of various types of content, from photography, art, and graphic design to audiovisual media to text.

On behalf of my firm, we fully agree with and support your AI Registration Guidance, which has "reiterated the principle that copyright protection in the United States requires human authorship." In our professional opinion, no work, be it audiovisual, graphic, design, text-based, music, or other media or formats, should be protected with copyright if it was created in whole or in part by algorithms or computer code, such as Large Language Models, Generative Pre-trained Transformers, or other forms of artificial intelligence or computer systems. In other words, if a non human entity or process is directed by either a human or another computer to create a work, that work should not be granted copyright, and the human individual (or corporate entity) who guided the creation of said work should never be considered the copyright holder. More specifically still, if the non human entity or process (for example an AI system) employed in the creation of said works, had been trained on human-produced works without proper notification, consent, credit, or compensation, it should be deemed that any outputs are at least partially or indirectly plagiarized and therefore do not qualify to be considered "original works" in any sense of the word.

Granting copyright to AI- or machine-generated works would not only grossly devalue and invalidate the talent, creativity, ingenuity, and effort of millions of human individual creators around the world, it would effectively put into question the right to their own works, and diminish their potential to earn a living using their own skills. This would in turn destabilize the U.S. economy and likely the economies of other nations as well.

We urge the U.S. Copyright Office to stand firmly in support of the millions of creators and entrepreneurs in the United States, and protect their right to own their own original works and protect the concept, practice, and legal status of human authorship.

Thank you for your consideration.

Comment from Creative Commons

Comment ID: COLC-2023-0006-8735

Response from Creative Commons

Submitted by Kat Walsh, General Counsel, on behalf of Creative Commons

Creative Commons is writing to address the questions posed in the Notice of Inquiry on Artificial Intelligence and Copyright, Docket 2023-6. This response addresses questions 1, 2, 4, 5, 8-10, 15, 18-22, and 32.

General Questions

(1,2):
Creative Commons is a nonprofit organization that helps overcome legal obstacles to the sharing of knowledge and creativity. Toward that end, we publish a suite of licenses and legal tools designed to give every person and organization in the world a free, simple, and standardized way to grant copyright permissions for creative and academic works; ensure proper attribution; and allow others to copy, distribute, and make use of those works. In addition, we work closely with major institutions and governments, including the United States government and the governments of numerous U.S. states, to create, adopt and implement open licensing and ensure the correct use of CC licenses and CC-licensed content.

A large portion of CC license users are artists and creators, particularly independent artists. These creators are choosing to give away some of the exclusive rights granted to them by default by applying CC licenses to their works, to best fulfill their own economic and creative purposes. We believe that generative AI presents both benefits and risks, but that there is tremendous potential for both the creative and scholarly communities.

Creative Commons has always sought out ways to harness new technology to serve the public interest and to support better sharing of creative content — sharing that is inclusive, just, equitable, reciprocal and sustainable. We support ways for creators to share their works as broadly and openly as they want, so that people can enjoy them globally without unnecessary barriers. We also advocate for policies that ensure new and existing creators are able to build on a shared commons, while respecting creators' legitimate interests in control and compensation for their creative expressions.

A founding insight of Creative Commons is that all creativity builds on the past. Creators necessarily learn from and train their own skills by engaging pre-existing works and artists — for instance, noticing the style in which musicians arrange notes, or building on surrealist styles initiated by visual artists. Star Wars invented the character of Luke Skywalker, but it built on the idea of the hero's journey, among many other elements from past works. People observe the ideas, styles, genres, and other tropes of past creativity, and use what they learn to create anew. No creativity happens in a vacuum, purely original and separate from what's come before.

Generative AI can function in a similar way. Just as people learn from past works, generative AI is trained on previous works, analyzing past materials in order to extract underlying ideas and other information in order to build new works. Image generation tools develop representations of what images are supposed to look like by examining pre-existing works, associating terms like "dog" or "table" with shapes and colors such that a text prompt of those terms can then output images.

Given how digital technologies function, training AI in this way necessarily involves making temporary copies of images in order to analyze them. We think this sort of copying can and should be permissible under copyright law. There are certainly nuances when it comes to copyright's interaction with these tools, some of which we will address further in this document. But treating copying to train AI as per se infringing copyright would in effect shrink the commons and impede others' creativity in an over-broad way. It would expand copyright to give certain creators a monopoly over ideas, genres, and other concepts not limited to a specific creative expression, as well as over new tools for creativity.

We have concerns about this technology and its responsible development and use: the impact these tools will have on artists and creators' jobs and compensation, use of these tools to develop harmful misinformation, to exploit people's privacy (eg, their biometric data), or perpetuate biases, as well as the sustainability of the information commons. These issues particularly impact the open sharing community, which often is producing information, scholarship, and creative output in the public interest, without the same incentives for creation as many other content creators. But copyright is not the only lens to address these problems, and our response in general cautions against using copyright as a tool to address all social concerns in place of more appropriate remedies.

Toward these ends, we have had several consultations with our communities on AI policy and approach with the aim of growing and sustaining the commons of freely-available information, some of which will be referenced below. We have heard from many different people, including artists and other rightsholders as well as developers of AI technologies, and we believe that it is possible for policy to guide the development of AI in ways that benefit all. These conversations inform our responses, here.

(4):

Creative Commons has been working with the EU to refine its approach to AI regulation, the AI Act. While that law is not perfect, its overall approach to AI regulation is logical insofar as it focuses on harmful uses rather than the technology itself. That said, this approach may have difficulty addressing advances in the technology because of the broad scope of potential uses that it holds; any approach to regulating AI must be flexible and able to develop with changes in the technology. We are also concerned about the scope and content of some of the requirements on disclosure of training data, which, while generally desirable, may be implemented with requirements that are too burdensome for practical compliance.

We have also seen Japan's approach making permission to train AI explicit: in July 2021, Japan's Ministry of Economy, Trade, and Industry published its report on AI Governance in Japan which provisionally disavowed any need to institute "legally-binding horizontal requirements for AI systems," while also stating that future discussions should "be implemented in consideration of not only risks but also potential benefits."[1] Keiko Nagaoka (Minister of Education, Culture, Sports, Science, and Technology) confirmed that existing Japanese law does not outlaw uses of copyrighted material for the training of generative AI. This policy clears up uncertainty and places Japan in a key position to be a center for AI development.

Additionally, CC believes that international consistency would be strongly desirable for AI regulation. Not only will this help individuals so they do not have to navigate a maze of regulations when trying to understand their rights in an individual jurisdiction, it can help to avoid a race to the bottom, where AI development concentrates in the place with the least regulation.

(5):
We don't believe new copyright legislation is warranted to address copyright-related issues with generative AI. In particular, we caution against the creation of new rights that would cover uses that currently do not require a license under limitations and exceptions to copyright.

We are concerned about new restrictions on the ability to make use of a copyrighted work in order to extract data from it to be processed by a machine. Treating copying to train AI as per se infringing copyright would in effect shrink the commons and impede others' creativity in an over-broad way. It would expand copyright to give certain creators a monopoly over ideas, genres, and other concepts not limited to a specific creative expression, as well as over new tools for creativity.

We also see the copyrightability of outputs as being adequately handled by existing copyright law: that works not authored by humans should not be eligible for copyright protection, and that works authored by humans should be eligible. There is a tremendous variety of tools and processes lumped together under "generative AI" which involve varying degrees of human involvement in the creative process, some of which the Copyright Office has already considered. We consider the approach of granting copyright to the human-authored parts of a work while exempting the AI-generated parts to be the correct approach, even where drawing the distinction in practice may be difficult. This problem is not unique to AI: authors have always created work building upon and remixing uncopyrightable and other public domain works, by remixing old songs, translating or editing old texts, incorporating government-authored works, and other such artistic components from the public domain.

1 https://www.meti.go.jp/shingikai/mono_info_service/ai_shakai_jisso/pdf/20210709_8.pdf.

(8, 8.1, 8.5):
We believe that, in general, use of copyrighted works for training should be fair use.

In particular, we believe the uses being made of these works are transformative, fitting in with a long line of cases that have found that non-consumptive, technological uses of creative works in ways that are unrelated to the expressive content of those works are transformative fair uses. In *Author's Guild v. Google*, the Second Circuit Court of Appeals found that Google's act of digitizing and storing copies of thousands of print books to create a text searchable database was fair use, as Google's purpose was different from the purpose of the original authors. The books were like pieces of data that were necessary to build Google's book database, and they were not being used for their expressive content. It was necessary to use the copyrighted material to build a digital tool that permitted new ways of using print books that would be impossible otherwise; the books as part of Google's database served a very different purpose from their original purpose, which supported the finding of fair use in this case.

It is also similar to the search engine case *Kelly v. Arriba Soft*. In this case, search engine Arriba Soft copied and displayed copies of photographs as thumbnails to users. The court held that this use served a different and transformative purpose from the original purpose because Arriba Soft only copied Kelly's photographs to enable its search engine to function, and not because of their aesthetic value. The images here served a function as data for the tool, not as works of art to be enjoyed as such.

The works being used as data to train these AI systems function similarly. Copyrighted works are used as sources of data from which data is collected and used to form mathematical models that learn from this data. These models, which have drawn inferences from the works but do not contain exact copies, can then generate new works of computer authorship which are also not copies of the original works.

Andy Warhol Foundation v. Goldsmith also continues to support this conclusion that training generative AI is transformative fair use. *Warhol* directs lower courts to compare the specific ways that an original work and a secondary work are used. In that case, the Court held that two commercial magazine uses of a photo had identical purposes, suggesting that the fair use analysis of this factor would be different if the uses had differed in character and purpose. With generative AI, the use of copyrighted works for AI training has a fundamentally different purpose from the aesthetic purposes of those works. This use bears more resemblance to the transformative use in search engines than the similar uses of artistic display of the subject of the image in commercially published works.

Notably, training generative AI is separate from using generative AI. While it may be possible to generate works to use generative AI in ways that infringe upon pre-existing works and that compete in the market for the works that the models are trained upon, the training itself does not interfere with the original intended aesthetic use or value of the originals. Accordingly, the market considerations should be considered only where the work is being used in a form that duplicates the original–simply gleaning data from a work is not sufficient, as the newly created work is not a comparable market substitute. As *Warhol* noted, the "central question" in fair use is whether a secondary use "supersedes the objects" of the original work. Using creative works to learn about what makes them what they are does not interfere with or supplant those works.

(9):

Creative Commons recognizes the value of giving creators ways to define their preferences for how others use their works; we have a suite of six separate licenses that millions of creators use to share their works with the commons according to their intentions for the work, built atop copyright law rather than replacing it.

We see value in approaches allowing rightsholders to affirmatively express preferences to opt in or opt out of training, particularly as addressing these preferences is likely to create more trust in the systems where the development has respected them. But we are skeptical of this as a new requirement in copyright law–instead, we would like to see voluntary schemes, similar to approaches to web scraping, which allow for standardized expression of these preferences without creating strict barriers to usage in cases where it may be appropriate–for example, use of this data in training AI models for academic research purposes.

For another example, CC is supportive of traditional knowledge labeling such as the scheme developed by Local Contexts.:[2] This scheme does not rely solely on rights in copyright law; instead, it calls for a voluntary standard as to how this material should be treated because copyright law is unsuited to the needs of these communities.

In our community consultations on this issue, we saw general support for this area of exploration but no clear consensus.[3] While many artists and content creators want clearer ways to signal their preferences for use of their works to train generative AI, their preferences vary. Between the poles of "all" and "nothing," there were gradations based on how generative AI was used specifically. For instance, they varied based on whether generative AI is used

- to edit a new creative work (similar to the way one might use Photoshop or another editing program to alter an image)
- to create content in the same category of the works it was trained on (i.e., using pictures to generate new pictures)
- to mimic a particular person or replace their work generally, or
- to mimic a particular person and replace their work to commercially pass themselves off as the artist (as opposed to doing a non-commercial homage, or a parody).

Views also varied based on who created and used the AI — whether done by researchers, nonprofits, or companies, for instance.

Many technology developers and users of AI systems also shared interest in defining better ways to respect creators' wishes. Put simply, if they could get a clear signal of the creators'

2 https://localcontexts.org/labels/traditional-knowledge-labels/
3 https://creativecommons.org/2023/08/31/exploring-preference-signals-for-ai-training/

intent with respect to AI training, then they would readily follow it. While they expressed concerns about over-broad requirements, the issue was not all-or-nothing.

While there was broad interest in better preference signals, there was no clear consensus on how to put them into practice, with concerns about social effects being expressed on both the pro and con sides of introducing it. The effect on the information commons is not yet clear; a hard requirement enshrined in law at this stage may lead to choices that do not serve the public interest. We are aware of numerous efforts from AI developers on systems that respect voluntary preference signals, and we would like to see further work and experimentation with voluntary schemes before supporting a legal requirement.

(10, 10.4):
CC is aware that Creative Commons licensing is often pointed to as an example of extending creator choice through licensing schemes, and as the steward of those licenses we caution against relying on that example. Primarily, the CC licenses are drafted explicitly to extend only as far as copyright, and not to override relevant limitations and exceptions; it does not create contractual obligations where a use would be allowable under fair use. That is, where a use of a copyrighted work would be permitted by fair use, the CC license does not impose its conditions, such as attribution requirements or commercial use restrictions.

We are skeptical of a collective licensing approach because of the importance of maintaining limitations and exceptions to copyright for de minimis and transformative uses.

Transparency & Recordkeeping

(15):
CC generally supports measures that provide the public with transparency into AI models and the data used to train them. Transparency can help build trust in AI models by allowing others to "look under the hood" and investigate how those models work. However, these measures should not be so onerous as to make it impractical or impossible for developers to develop their technology. Moreover, transparency should not be enforced through existing or newly created copyright measures.

We believe disclosure of training data in the development of AI models is desirable for numerous reasons, including suitability for academic and research purposes, supporting the sustainability of communities and institutions where training data is created, and crediting sources. However, we do not support an obligation for disclosure of training data that is rooted merely in copyright.

For data sets of copyrighted works, we encourage collection and disclosure of licensing and authorship information as far as is possible when it is available for those works and would expect best efforts to be made to retain them. At the same time, we recognize there may be some important data sets or use cases where such information is not reasonably possible to collect or distribute.

Copyrightability

(18-21):
Human creativity should be eligible for copyright, even where combined with AI-generated work. The current standards for copyrightability of a creative output should remain as they are in the context of AI, where human authorship is eligible for copyright and machine authorship is not. We recognize that determining which parts of a work are authored by a human in such situations will not always be clear–indeed, as determining authorship of individual parts of a work is not always clear in works created without AI assistance.

Copyright protection for the code should provide sufficient incentives to develop these systems without relying on additional rights in the outputs. Indeed, granting developers rights in the outputs would both stifle creative uses of these tools and be inconsistent with previous precedent. It is true that the technology shapes the artistic output of the tool–as does every tool for creating artwork. Tools for creating increasingly sophisticated artworks have been treated as just that, tools, and merely the implements that help creators bring forth their artistic visions. The Supreme Court looked at the relationship between artists and their mechanical tools vis-a-vis copyright over 100 years ago in *Burrow-Giles Lithographic Co. v. Sarony* and found that works produced with the help of artistic tools only qualify for protection insofar as the artists who use them add their own personal creativity to their outputs. We believe that works produced with the assistance of generative AI should be treated in a similar way.

Infringement

(22):
It is possible for AI-generated works to implicate the exclusive rights of existing rightsholders in some cases. We cannot enumerate all cases specifically; however, there are many possible ways this can happen, and we urge caution and flexibility in determining when an output has infringed exclusive rights. It is possible for an output to be infringing; it is also possible that an output including material under copyright may be using it in a way that is a fair use.

Liability for any infringement in this case should depend heavily on the facts of the situation, as the nature of generative AI tools and usage of those tools can vary greatly. For example, where a user has made a knowing attempt to get copyrighted material they have previously had access to as output from an AI tool (by directing prompts specifically toward that output, for example), that user has simply used the tool to reproduce the material, in the same way that one might use other tools to create a copy of an existing work.

We would caution against considering a simple attempt to make reference to a rightsholder's work as an infringement—for example, a user requesting artwork in the style of an existing artist. Copyright law already contemplates such cases, and already deals with some ambiguity that depends on facts: was a creator simply being inspired by the style of an artist, which is not an infringement and is critical to free expression and the development of artistic genres, or attempting to reproduce copyrightable elements in an infringing manner?

The debate about how copyright should apply to generative AI has often been cast in all-or-nothing terms—does it infringe on pre-existing copyrights or not? The answer to this question is certainly that generative AI *can* infringe on other works, but just as easily it may not. The tools that already exist in copyright, like the tests for substantial similarity and fair use, are sufficient to address these questions. Our concern is that overregulation of generative AI could easily destroy this powerful tool that enables new ways for people to express their creative visions and share their works with the world.

Other questions

(32):
CC agrees with long-standing doctrine that style is not protected by copyright. As the Supreme Court recently recognized in *Andy Warhol Foundation v. Goldsmith*, the idea-expression doctrine provides a safeguard against copyright encroaching on First Amendment rights. If copyright begins to recognize protection over style because of generative AI, it runs the risk of granting rights to ideas in ways that are contrary to the fundamental principles of copyright law and that can harm the ability of future creators to express their artistic visions.

Conclusion

These questions about how copyright law should apply to generative AI are often difficult, fact-specific, and touch on areas other than copyright as means to solution. Copyright is an appealing road to look down for numerous reasons, particularly as it comes with easily available remedies and contractual alterations, but many of the questions presented imply more difficult questions about labor, ethics, misinformation, and other areas which the Copyright Office is not the correct place to address. As such, we suggest caution and flexibility in addressing generative AI. As with many new technologies, it presents great promise along with its negative effects, and we are concerned about addressing those effects with requirements that do not adequately consider the tradeoffs.

Creative Commons has seen its licenses and legal tools used as the legal infrastructure for numerous websites and millions upon millions of creative works by rightsholders who see the value of open sharing to a more creative, more knowledgeable, more equitable, and more innovative society. We have advocated for copyright as a balance between the individual rightsholder and the public interest. Limitations and exceptions to copyright form an essential part of US copyright law in contemplation of that public interest, and in this response we ask that those limitations and exceptions be considered as the essential feature that they are rather than simply considering how best to assign the exclusive rights.

Comment from Fiction Factory Games LLC

Comment ID: COLC-2023-0006-3016

My name is Stefan Gagne, owner of Fiction Factory Games LLC, and a citizen of Maryland.

I work in an industry potentially under threat by A.I. generated content -- video game development. The threat is twofold; the use of cheap A.I. generated content to replace artists, writers, programmers, and musicians. Also the threat of our own art, writing, programming, and music (all of which is copyrighted) being used to train an A.I. which can then generate content similar to ours and sell it cheaply.

A.I. models are being trained on copyrighted material regularly, with no opt-in, no choice, no consent given to have our work repurposed and remixed and re-sold commercially. Because millions of creators are having their work stolen in this manner it becomes difficult if not impossible to point to any one portion of an A.I. generated work and say "that's copying my work." This is not similar to how a human takes inspiration from another's work -- most A.I. models literally could not exist without stolen assets.

I don't believe A.I. technology is innately problematic. If trained on a legal set of materials, such as writing that has been commissioned expressly for purposes of training the model, then it can be a very powerful and beneficial tool. But the current practice of "scraping" the web for anything publicly viewable is wildly unethical -- just because words or images are on the internet doesn't mean they're free from copyright, a fact these systems completely ignore, vacuuming them all up into their models without attribution and without checking legality.

Regarding the issue of whether or not A.I. generated content can be copyrighted, I don't feel it can be unless the model is proven to use 100% ethically sourced and copyright-safe materials. Without that, the resulting output invariably is using illegally obtained material and cannot be considered a derivative work as no intentional and consensual human creativity went into its making. The issue is the obvious loophole -- if A.I. content is not copyrightable, all someone has to do is generate content, then adjust it very slightly by changing a word here and there or doing minor editing, and now it's arguably copyrightable. This trivial effort to skirt around the law is difficult to deal with and honestly I don't have a solution, as I am not a lawyer. But I hope for the sake of millions of artists who are harmed by unrestrained use of A.I. that an answer can be found.

I'd like to close out this comment with a warning -- because A.I. is capable of generation hundreds, thousands, even millions of similar but unique texts, there's a strong worry this very comment submission process will be flooded by A.I. advocates using A.I. itself to create pro-A.I. submissions under fake names. I hope you're prepared for the potential tidal wave of misinformation that will come with opening this public inquiry.

Please do the right thing. Rule against the violation of copyright inherent in most A.I. models. As human artists we do not consent to have our work repurposed and copyright-claimed by others. Please stop this before it results in the collapse of our nation's creative enterprises. Thank you for your time.

Comment from Black Tabby Games

Comment ID: COLC-2023-0006-6780

As the founder of an award-winning video game studio, I have a large body of concerns revolving around intellectual property rights as they pertain to generative AI models. And before I outline those points, I would also like to stress that the bulk of these concerns are held by practically every professional artist and creator I know, particularly those who own their own intellectual property.

1. Generative models are not true artificial intelligence — they are probability based algorithms — and by extension, the work they perform is incapable of being transformative.
2. Generative AI models rely on mass amounts of training data obtained without consent, and without the opportunity for artists to opt or, or, in cases where they would opt in, without any mechanism to compensate said artists for the use of their labor. The most mainstream models, including Stable Diffusion and Midjourney, have scraped most images from the internet as part of their training data without consent.
3. Though proponents of AI Gen models will argue that diffusion-type algorithms generate unique images, repeated examples of generated images veering remarkably close to training data—including artifacts the retain part of an artist's signature—show that this is not the case.
4. The human-level input on AI Gen image and text generation is minimal — it is a series of single word or short phrase prompts that spit out dozens of images. If we agree, for instance, that by default, the intellectual property rights for a commissioned piece of art belong to the artist rather than the commissioner, there is no reason to believe that using an AI-gen model is any more transformative than providing a description to a human artist, aside from the fact that the AI's work itself is not transformative.
5. Since nothing transformative takes place, and since the model relies on billions of points of data without the license to use it, I believe that beyond not being copyright-able, all AI art generated by those algorithms is infringing on some level.

I don't have copyright related concerns with AI-gen models where the user owns a license for all of that model's training data, but given the massive scope of data necessary for a model to function, I find it hard to believe that those hypothetical models would be functional.

I have a number of other issues that extend beyond copyright implications, but given the focus of your office, I will hold those concerns for this comment. The negative implications of these tools in terms of disrupting the labor market and negatively affecting the arts, however, is astronomical.

I can also say that has someone who travels in the professional circles of artists and other creators, my concerns are ubiquitous across that section of the economy. Absolutely nobody, aside from self proclaimed "disruptors" has a favorable view of artificial intelligence on the arts, and all share concerns around their work being used without their consent on an industrial scale to try and replace them.

Thank you for taking the time to read my comment,
Tony Howard-Arias (President, Black Tabby Games)

Comment from Ace Of Words

Comment ID: COLC-2023-0006-0729

Language models, art generators, and the (currently experimental) video generators which are named "AI" are trained by exposure to data sets. In order for an AI to produce coherent, cogent, and lifelike results, those data sets must be language OR artwork OR video produced by real humans. In the case of most of the dominant "AI" available on the internet today, these data sets were produced by "scraping," the process of moving from one webpage to the next in search of publicly available information--I stress, publicly available, but NOT COPYRIGHT FREE.

As a result, it seems apparent right out of the gate that these "Artificial Intelligence" models cannot produce material that can be copyrighted. Just because material is publicly available, does not mean its owners no longer have the right to it. And "AI" is not like an artist. It cannot produce an original work. If it produces a drawing of an orange, that drawing is inexorably and inextricably linked to all the other (likely copyrighted) pieces of art which the AI learned from.

As a consumer of art, and a producer of writing and digital artwork, I cannot overstate how strongly I oppose even the idea of "AI" generated content being allowed anything even approaching copyright, because "Artificial Intelligence" as it now stands cannot create--it can only reproduce what it has seen.

Anything an AI makes cannot be copyrighted. Anything made by an AI ought to be ipso facto non-restricted public domain unless modified significantly (arbitrary number, modified by 50% or more) by a human. Further, there must be some way for humans to "opt out" of inclusion in AI training data sets. There should be significant financial and potential civil penalties for companies who trawl the web indiscriminately to rob from writers and artists around the world. Thank you for your time, and please make the right decision to ensure that future generations will be able not only to create, but to envision a life as creators.

Pearson

Electronic submission via www.regulations.gov

Comment from Pearson

Comment ID: COLC-2023-0006-8703

October 30, 2023

Artificial Intelligence and Copyright (Docket No. 2023-6) Notice of Inquiry and Request for Comments

Introduction

For many decades, Pearson has been a leader in educational publishing. Today, we are the world's leading learning company, with more than 160 million users of our products and services worldwide. Approximately 20,000 Pearson employees, together with our authors, are committed to creating enriching learning experiences designed for real-life impact. We are invested in the future of learning, from its ability to unlock opportunities for learners and societies globally to learning through innovation. Learning is who we are.

Pearson is a trusted brand that represents high quality product offerings and enriching, vibrant learning experiences. Our company's unparalleled collection of high quality, proprietary intellectual property assets strengthen our ability to deliver an unmatched experience for the learner across their lifetime of learning. With our unrivalled depth of content and data, we expect generative AI to create significant positive opportunities for Pearson. However, like all new technologies, the use of generative AI must be approached in a thoughtful and considered manner to avoid doing more harm than good. As learners and educators place enormous trust in Pearson, it is essential that our content and copyrights be protected. Pearson welcomes the opportunity to participate in the United States Copyright Office's notice of inquiry on generative AI. While there are other risks associated with the use of generative AI, including cybersecurity, data privacy, bias and unfairness, our responses below focus on and illustrate the importance of copyright in our protection of our valuable business assets.

Pearson's Business

Pearson's purpose is to provide life to a lifetime of learning. We do this by creating vibrant learning experiences for learners in many settings including primary and secondary schools, higher education, the workforce, professional assessments, and English language programs. We create reliable and up-to-date materials by collaborating with trusted authors and expert content creators such as teachers, faculty, and educational institutions. The high-quality learning experiences we provide include not only literary and artistic works, traditionally in written text, but also data in many forms, including digital books, digital learning environments, audio visual content, computer programs, databases, maps, and drawings. This effort takes considerable investment in both time and resources and results in highly valuable creative products that are protected by copyright.

The worldwide educational publishing market was estimated to be worth more than 6 trillion U.S. dollars in 2022 and is expected to grow to 8 trillion U.S. dollars by 2030.[1] It contributes hundreds of millions of dollars to the U.S. economy and supports millions of high value jobs. It is essential that this socially important and valuable industry enjoys the protections under the law to which it is entitled.

General Questions

1. As described above, generative AI systems have the ability to produce material that would be copyrightable if it were created by a human author. What are your views on the potential benefits and risks of this technology? How is the use of this technology currently affecting or likely to affect creators, copyright owners, technology developers, researchers, and the public?

We are cautiously optimistic about the potential of generative AI to improve education and learning, provided the appropriate steps are taken to mitigate the risks. We see potential benefits in personalized learning, increasing access and equity, improving outcomes, and in feedback and engagement. However, the risks of generative AI include cybersecurity, data privacy and the introduction of bias and unfairness, as well the delivery of faulty, unreliable data.

Unfortunately, the lack of transparency about generative AI training data and processes makes it difficult to know if a specific generative AI system complies with all relevant copyright law, other intellectual property laws, database rights, data protection laws, and international laws and treaties, such as the Berne Convention. It is essential that generative AI companies be transparent about the content and material used to train their systems, trace the provenance of such content by disclosing their processes, track the copyrighted works used, and retain the copyright management information of such content.

With regard to copyright interests in particular, most generative AI systems currently available have been trained using copyrighted content without the authorization of rights holders. These impermissible activities must be addressed not only because they use valuable content without the remuneration to copyright holders, but also because the unauthorized and unrestrained use of copyrighted materials disincentivizes innovation, discourages creation of new high-quality content, and results in inaccurate and less valuable machine-produced derivative content and works.

Fundamentally, the exclusive rights of copyright should not be ignored, and generative AI companies should not be allowed to avoid or circumvent the copyright or other intellectual property rights of others simply in the name of innovation. Copyright holders have in no way automatically relinquished any exclusive rights in

[1] *Global Education's $8 Trillion Reboot*, MORGAN Stanley (June 7, 2023), https://www.morganstanley.com/ideas/education-system-technology-reboot ("The global education market is slated to reach an estimated $8 trillion in value by 2030 from $6 trillion in 2022.").

their content or given up their right to consider market or licensing opportunities for their content in the generative AI context. It is critical that generative AI companies obtain the consent of the copyright owners for the use of their content to avoid infringement and the other negative outcomes discussed above. Pearson's content is much more than just the 1s and 0s of "data", it represents valuable creative assets developed by experts with substantial investment that generative AI developers seek to use to produce quality models and outputs.

2. Does the increasing use or distribution of AI-generated material raise any unique issues for your sector or industry as compared to other copyright stakeholders?

As the world's largest learning company, we see a variety of unique issues related to the increasing use or distribution of AI-generated material in the education and workforce space.

First, the use of AI systems for cheating and plagiarism is well-known. The availability of these tools for unethical purposes poses challenges not only for students, teachers, and the integrity of academic systems, but also undermines the usefulness and market for educational products. If AI companies negotiated appropriate licenses with educational publishers, we might be able to set limits on the amount of any one work that is ingested and the sort of output that the program could deliver, making AI systems less prone to serve as tools for cheating.[2]

Second, educational publishers like Pearson are particularly concerned with the reliability and validity of outputs generated by these generative AI systems. In the learning environment, high-quality, vetted content is critical to providing learners with positive outcomes. Likewise, accuracy and pedagogy are of paramount value. We understand that generative AI systems have used material they have sourced indiscriminately from the internet to train their models. The quality of this material, its accuracy, and its sources are unverified, and these systems return outputs that are inaccurate or misleading. To add to our concern, we have seen generative AI models produce inaccurate content and misinformation and attributed that content to Pearson products.

That generative AI models output incorrect responses and attribute them to authors and publishers who have spent significant resources in establishing their favorable goodwill and strong reputations in delivering quality educational material is alarming. It will remain critical in the learning industry that humans are involved in the creation, development, and deployment of generative AI systems to ensure responsible sourcing of training data and use of these technologies.

Third, we believe the scale of unauthorized copying of educational materials is likely greater than other forms of content. As discussed above, developers of generative AI systems use datasets comprised of data that has been indiscriminately sourced from the internet. They use automated web-crawling tools to ingest large volumes of data, including huge datasets of eBooks, often pirated, such as the "Books3" repository and

[2] To be clear, the opportunity to control the use or non-use of an entity's intellectual property, including whether or not to license its intellectual property, vests with the intellectual property owner.

others.[3] The unlawful presence of pirated material on the internet is a battle that educational publishers continue to fight, but the unchecked sourcing and scraping of educational content and material from around the internet to train generative AI models further propagates the harms of online piracy by dramatically expanding the reach of these pirated sources.

Fourth, the use of generative AI risks undermining copyright's core value of promoting an incentive to create.[4] It is no stretch to believe that AI systems have been or will be used to create wholly AI-generated textbooks trained on and designed to replace those of responsible publishers like Pearson. Not only would such AI-generated content infringe rights in the underlying content on which such outputs produced and risk publication of unreliable content, but it would materially impact the incentive of true creators to create new works for the advancement of the arts and sciences.

4. Are there any statutory or regulatory approaches that have been adopted or are under consideration in other countries that relate to copyright and AI that should be considered or avoided in the United States? How important a factor is international consistency in this area across borders?

The United States should carefully consider other countries' approaches to copyright and AI and the implications those approaches have in the U.S. Of note are two international examples below.

European Union. The European Parliament's proposed EU AI Act would require providers of generative "AI Systems" to publish summaries of copyrighted data used for training the AI's model.[5] Specifically, providers of generative AI systems would need to "make publicly available a sufficiently detailed summary of the use of training data protected under copyright law."[6] Generative AI companies in the United States should face the same requirement.

Requiring transparency and adequate disclosures of training data from the developers of generative AI models is necessary for several reasons. First, it would help copyright holders protect their works. If generative AI models were required to provide a detailed summary of their training data, this would better

---

[3] Alex Reisner, *What I Found in a Database Meta Uses to Train Generative AI*, THE ATLANTIC (Sept. 25, 2023), https://www.theatlantic.com/technology/archive/2023/09/books3-ai-training-meta-copyright-infringement-lawsuit/675411/.

[4] *See, e.g.*, David De Cremer et al., *How Generative AI Could Disrupt Creative Work*, Harv. Bus. Rev. (Apr. 13, 2023), https://hbr.org/2023/04/how-generative-ai-could-disrupt-creative-work ("[G]enerative AI significantly changes the incentive structure for creators, and raises risks for businesses and society. If cheaply made generative AI undercuts authentic human content, there's a real risk that innovation will slow down over time as humans make less and less new art and content.").

[5] Eur. Parliament, *EU AI Act: first regulation on artificial intelligence* (June 14, 2023), https://www.europarl.europa.eu/news/en/headlines/society/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence. The European Union is currently considering the EU AI Act. The EU AI Act would cover any "AI system" that is developed by an EU-based provider, as well as any systems that are developed outside of the EU and placed on the EU market. *Proposal for a Regulation of the European Parliament and of the Council Laying Down Harmonized Rules of Artificial Intelligence (Artificial Intelligence Act) and Amending Certain Union Legislative Acts*, COM (2021) 206 final (Apr. 21, 2021). However, the member states are still negotiating the definition of AI System. It remains to be seen whether there will be consensus or alignment on a singular definition or whether individual governments will be left to define AI Systems independently during transposition, potentially creating inconsistency in the Act's implementation.

[6] Eur. Parliament, *EU AI Act: first regulation on artificial intelligence* (June 14, 2023), https://www.europarl.europa.eu/news/en/headlines/society/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence.

facilitate a determination of whether copyrighted works have been used without authorization. Second, mandatory disclosure would likely help to produce AI models that are of superior quality. Currently, serious allegations have been made against generative AI models for sourcing third-party copyrighted materials from pirated websites.[7] In addition to infringement, this type of illicit sourcing contributes to poor quality training data and reflects poorly on the rights holders of the copyrighted material that has been collected from unauthorized sources. Consistent, mandatory disclosure requirements would disrupt this unauthorized sourcing and help ensure that AI models train on high-quality, authentic, and authorized works. Consistent disclosure requirements and transparency surrounding the use of materials used in the training process would also improve the ability of users and government entities to scrutinize the training process. Importantly, such requirements likely would spur self-regulation by the AI companies to avoid running afoul of copyright law in the first place.

Under European Union copyright law, there are two exceptions for content scraping (sometimes referred to as text and data mining ("TDM")), including an exception for research and cultural heritage institutions and an exception that covers TDM for any purpose.[8] This second exception grants copyright holders the right to opt out of the exception and expressly reserve such use of works to themselves. While this approach reflects a compromise between the technology sector and rights holders, it does not, in our view, fully protect the rights holders. Rather, it places an impossible burden on rights holders to notify each and every AI company of its intent to opt out of content scraping. Essentially, it creates a new "formality" on rights holders and inverts the benefits of ownership. Copyright ownership does not require that the owner of that right affirmatively declare that a third party does not have permission to use the owner's copyrighted work. Instead, copyright ownership gives that rights holder the right to control use of their work - to decide whether or not to grant someone permission to use that work and whether to keep that work for their own use. No affirmative act by the copyright owner such as an "opt out" is currently required. A better approach would be analogous to Australia's copyright law, described below, which requires a generative AI company to affirmatively seek the permission of a copyright holder to use its content.

Australia. Under Australian copyright law, there is no exception for TDM for AI development. Accordingly, in Australia, to engage in TDM to train an AI model, an AI company must affirmatively seek and obtain the permission of a copyright holder. This approach recognizes the priority, import and breadth that should be placed on copyright, and it ensures that copyrighted materials are more fully protected, without putting the onus on the copyright holder to continue to justify and establish its rights each time an AI model is developed. This Australian law withstood lobbying efforts and claims that the approach was inflexible and would significantly damage and even paralyze innovation in Australia compared to other parts of the world. Notwithstanding this resistance, AI innovation does not appear to have been foreclosed or even been stifled in Australia. Thus, such concerns should be considered with extreme caution as generative AI models

[7] *See, e.g.*, Complaint, *Authors Guild v. OpenAI*, No. 1:23 cv 08292 (S.D.N.Y. Sep. 18, 2023); Complaint, *Chabon v. OpenAI*, No. 3:23 cv 04625 (N.D. Cal. Sept. 8, 2023); Complaint, *Silverman v. OpenAI*, No. 3:23 cv 03416 (N.D. Cal. July 7, 2023); Complaint, *Tremblay v. OpenAI*, No. 3:23 cv 03223 (N.D. Cal June 28, 2023).

[8] EU Directive 2019/790 (April 17, 2019).

continue to forge ahead with development and deployment, even in jurisdictions that affirmatively require AI companies to obtain consent to ingest copyrighted materials.

8. Under what circumstances would the unauthorized use of copyrighted works to train AI models constitute fair use? Please discuss any case law you believe relevant to this question.

8.1. In light of the Supreme Court's recent decisions in Google v. Oracle America and Andy Warhol Foundation v. Goldsmith, how should the "purpose and character" of the use of copyrighted works to train an AI model be evaluated? What is the relevant use to be analyzed? Do different stages of training, such as pre-training and fine-tuning, raise different considerations under the first fair use factor?

8.2. How should the analysis apply to entities that collect and distribute copyrighted material for training but may not themselves engage in the training?

8.3. The use of copyrighted materials in a training dataset or to train generative AI models may be done for noncommercial or research purposes. How should the fair use analysis apply if AI models or datasets are later adapted for use of a commercial nature? Does it make a difference if funding for these noncommercial or research uses is provided by for-profit developers of AI systems?

8.4. What quantity of training materials do developers of generative AI models use for training? Does the volume of material used to train an AI model affect the fair use analysis? If so, how?

8.5. Under the fourth factor of the fair use analysis, how should the effect on the potential market for or value of a copyrighted work used to train an AI model be measured? Should the inquiry be whether the outputs of the AI system incorporating the model compete with a particular copyrighted work, the body of works of the same author, or the market for that general class of works?

While this inquiry asks about circumstances where the unauthorized use of copyrighted works to train AI models constitutes fair use, Pearson cannot point to any such examples with respect to its own works in the current AI landscape. Indeed, Pearson's position is that its copyrighted works cannot be used in AI training without our permission. Although fair use is a fact-specific inquiry, Pearson believes the case law fully supports this position, including the Supreme Court's recent decision in *Andy Warhol Found. for the Visual Arts, Inc. v. Goldsmith*, 143 S. Ct. 1258, 1275 (2023). In *Warhol*, among other things, the Court warned that "an overbroad concept of transformative use, one that includes any further purpose, or any different character, would narrow the copyright owner's exclusive right to create derivative works."[9]

A lack of fair use in the AI training context is also supported by the fact that entire copyrighted works, if not entire pirate databases consisting of tens of thousands of entire copyrighted works, have been ingested into generative AI systems without permission. And still further support for a lack of fair use can be found in the commercial nature of most generative AI systems, the absence of a fundamentally different purpose or use from the educational purpose for which educational works are created, and the market harm from lost sales, lost licensing revenue, and the usurpation of the right to decide whether to exercise a market

[9] 143 S. Ct. 1258, 1275 (2023).

opportunity or not. While by no means an exhaustive list of concerns, these issues warrant careful attention by the Copyright Office to the assertions of AI developers and their supporters with respect to fair use, especially those that advocate for any type of blanket fair use exception in the AI training context, which is wholly unjustified.

11. What legal, technical or practical issues might there be with respect to obtaining appropriate licenses for training? Who, if anyone, should be responsible for securing them (for example when the curator of a training dataset, the developer who trains an AI model, and the company employing that model in an AI system are different entities and may have different commercial or noncommercial roles)?

As a publisher, Pearson works every day to negotiate and obtain appropriate rights to publish content. When it comes to generative AI development, this process should be no different. If a company developing a generative AI product wants to use Pearson's copyrighted works in its training data, just like any other licensing prospect, Pearson's rights include the right to consider whether Pearson has an interest in licensing such content and whether mutually agreeable terms can be reached. The companies developing generative AI are some of the largest and most profitable in the world. They should not be given free rein to ignore the exclusive rights that have been conferred on copyright owners to pursue and promote their creativity. Further, allowing generative AI developers to shift the licensing burden onto rights holders would allow developers to gain an unfair advantage at the expense of rights holders and other stakeholders and undermine the rights vested in those who create, which are founded in the Constitution and enumerated in the Copyright Act.

18. Under copyright law, are there circumstances when a human using a generative AI system should be considered the "author" of material produced by the system? If so, what factors are relevant to that determination? For example, is selecting what material an AI model is trained on and/or providing an iterative series of text commands or prompts sufficient to claim authorship of the resulting output?

The capabilities of generative AI are just beginning to be understood. Given that the AI landscape is still being charted, we believe it is premature to foreclose the prospect that a human using a generative AI system could be considered the "author" of material produced by that system in limited circumstances. The formulation of a non-exhaustive list of factors would be useful to assessing whether there is sufficient human contribution and creativity to a generative AI output for that output to be protected by copyright. Moreover, the focus should not just be on the output. A similar list to assess how input, which is the product of sufficient human creativity, can establish or support copyright protection in certain outputs from a generative AI system would also be useful.

Generally, we would not expect that the act of providing an iterative series of text commands or prompts is sufficient to claim authorship of the resulting output. However, just as the Copyright Office suggests that copyright can only protect material that is the product of human activity, we believe that further consideration should be given to whether a claim of authorship in output may exist where the input itself is a representation of the original intellectual conception of an author and, thus, subject to copyright protection. For example, when a generative AI system is used to translate or copy edit a copyrighted work, is that resulting output sufficient for copyright protection because the input is protected notwithstanding that the translation or copy editing is provided by a generative AI system? We believe that these examples tend to illustrate the potential need for flexibility and further consideration to be given in the context of determining what human activity is needed for copyright protection in a work that uses generative AI, so that copyright holders do not experience unintended consequences from too stringent an application of the human component.

22. Can AI-generated outputs implicate the exclusive rights of preexisting copyrighted works, such as the right of reproduction or the derivative work right? If so, in what circumstances?

AI-generated outputs can certainly implicate the exclusive rights of preexisting copyrighted works, such as the right of reproduction or the derivative work right. As a threshold matter, Pearson notes that this question does not ask whether copyright infringement can arise from inputs but rather focuses solely on outputs.

Copyright holders enjoy exclusive rights in copyrighted works, including the rights to reproduce, distribute, perform, and display copies of the copyrighted work and to prepare derivative works based upon the copyrighted work.[10] These rights do not change simply because the potentially infringing works were partially or wholly created by software. Instead, the *prime facie* question of infringement focuses on the facial content of the allegedly infringing work. If, for example, an AI model outputted material substantially similar to a Pearson textbook, there could be no doubt that such output would implicate (and violate) Pearson's exclusive rights in copyright. Similarly, if an AI model were asked to generate a sequel to a copyrighted work, and that sequel would be a derivative work had it been created by a human given the same instruction, then that output would implicate (and violate) the copyright holder's right to prepare derivative works.

In short, there is no reason that an otherwise infringing output would fail to be considered infringing simply because it is AI-generated, nor is there any reason to allow AI developers to escape accountability in this situation.

26. If a generative AI system is trained on copyrighted works containing copyright management information, how does 17 U.S.C. 1202(b) apply to the treatment of that information in outputs of the system?

Under 17 U.S.C. § 1202(c), "copyright management information" ("CMI") is information conveyed in connection with copies of a work (including in digital form) such as its title, its author, its owner, terms and conditions for its use, and other information. This definition includes, for example, an author's name appearing next to copyrighted images.

For good reason, the statute protects against the removal or alteration of CMI. It is a violation of § 1202(b) to, with reason to know that it will induce, facilitate, or conceal copyright infringement, remove, or alter the information from a work, distribute it knowing that it was altered without the copyright owner's authorization, or distribute copyrighted works knowing that the information has been removed or altered without the copyright owner's authorization.

---

[10] *See* 17 U.S.C. § 106.

As with other aspects of copyright law, the answer here is not materially transformed in the generative AI context. If a person were to read an article that included CMI and then impermissibly distribute copies of the article with that information removed, that conduct would violate § 1202(b). Likewise, if an AI model removed CMI in its ingestion of a copyrighted work of a third party or were trained on that work with CMI and then distributed a copy of the work with the CMI removed, such conduct would violate § 1202(b).

Thank you for considering our views.

illum1na

Comment from Illumina, Inc.

Comment ID: COLC-2023-0006-10315

Re: Notice of Inquiry and Request for Comment Regarding Artificial Intelligence and Copyright Docket No. COLC-2023-0006

Illumina is a global leader in genomic sequencing and in the development of AI systems to accurately and efficiently interpret genomic data. We are grateful for the opportunity to provide feedback to the U.S. Copyright Office on copyright law and policy issues raised by artificial intelligence ("AI") systems.

We believe that AI has the power to transform the genomics industry, and to rapidly accelerate Illumina's mission of improving human health by unlocking the power of the genome. For example, we are already using AI systems today to identify, classify, and annotate genetic variants for a multitude of useful purposes. Our AI systems have helped researchers discover novel drug targets, identify clinically relevant genes, and uncover genetic drivers of rare and undiagnosed diseases.

We acknowledge that not all AI is the same, and the Copyright Office should consider the distinctions between different types of AI models. On one hand, there are large, foundational Generative AI models, trained on large swaths of the internet. On the other hand are more focused models: either Generative AI models that have been fine-tuned for specific domains, or non-Generative AI models tailored to their specific use cases. Our comments here are largely aimed toward the latter - focused, domain-specific models ("Focused AI Models").

The Copyrightability of AI Models and their Outputs

Copyright law, at its core, protects the rights of individuals to their original works of authorship - rights which are enshrined in the Constitution. Any sufficiently *creative* work of human authorship is suitable for copyright protection.[1]

We believe that, in many cases, the development of Focused AI Models undeniably satisfies the human authorship requirement *and* is more than sufficiently creative to meet the established threshold for copyright eligibility. We agree with the Copyright Office that these determinations will necessarily be fact specific and made on a case-by-case basis.

As noted in the *Copyright Registration Guidance: Works Containing Material Generated by Artificial Intelligence* (March 16, 2023)("2023 Copyright Office Guidance"), the US Supreme Court has defined an author as "[one] to whom anything ***owes its origin,*** originator, maker; ... "[2] It is our position that, in many instances, both Focused AI Models themselves and the output or material produced *via* Focused AI Models can and should be considered a creative work that originates with, or "owes its origin" to, a human. In other words, the human development and input into the AI model itself can under the appropriate circumstances be the "creative input" and/or "intervention from a human author," similar to the manner in which the choices a photographer makes supports human authorship in the output of a photograph.[3] However, we also believe that the threshold and factors relevant to determining "human

[1] Feist Publication Inc. v. Rural Telephone Service Co., Inc. 499 U.S. 340, 111 S.Ct. 1282, 113 L.Ed.2d 358 (S. Ct. 1991)
[2] *Burrow-Giles Lithographic Co.* v. *Sarony,* 111 US 53 (S.Ct. 1884) (holding a photograph copyrightable despite use of a camera to give "visible expression" to "ideas in the mind of the author").

Illumina • 5200 Illumina Way • San Diego, CA 92122 • Tel 858.202.4500 • Fax 858.202.4766 • www.illumina.com

[3] U.S. Copyright Office, Compendium of U.S. Copyright Office Practices sec. 313.2 (3d ed. 2021) ("Compendium (Third)")(stating "to qualify as a work of 'authorship' a work must be created by a human being" and that it "will not register works produced by a machine or

authorship" should be determined through case law in a manner similar to the threshold for creativity and other copyright principles.

Appropriate copyright protection for Focused AI Models and the original outputs produced via assistance from Focused AI Models will incentivize continued development of this technology and the creation of valuable works. Without these protections, authors may not invest the money, time, and/or energy necessary to develop new works.

We also believe that, in addition to incentivizing development, strong copyright protection will encourage Focused AI Model developers to make their models more broadly available. For example, a known feature of many AI models is that, with enough resources and effort, models can be used to reverse engineer their own training data. Without strong copyright protections for both independently created AI models *and* their original outputs, model creators would be limited to contract law remedies to protect their work, and would tend toward closed access to reduce the risk of misappropriation by third parties.

The Copyright Office's Approach to AI

Because the same principles of human authorship and creativity should apply equally to both AI and non-AI works, we believe that the Copyright Office's approach to AI should be similar to that of other works. Unfortunately, however, we find the current Copyright Office Guidelines for works produced by AI technology to be largely inconsistent with historical Copyright Office practices. For example:

- On the issue of originality, the Copyright Office historically has not required applicants to provide a detailed explanation of the independent creative process used by a named author to produce a work of authorship.
- On the issue of authorship, the Copyright Office historically has not required a photographer to identify each of the choices made when taking a photograph to satisfy the minimum threshold of human authorship.

In contrast, the current Copyright Office requirement that applications involving AI technology must describe detailed facts supporting human authorship in the material produced via the AI system[4] is inconsistent with historical Copyright Office practice and sets too high of a bar for registration.

It is important to note that our position regarding the current Copyright Office policies does not contradict well-established legal principles that the core of copyrightability is *human creativity*. However, the issue of precisely *how much* human creativity must occur, and at what point in time the human creativity must occur, is highly fact specific and should be determined on a case-by-case basis by the courts. For this reason, we encourage the Copyright Office to reconsider the current requirement that applicants must disclose their creative steps in order to satisfy the human authorship requirement. We also encourage the Copyright Office to avoid any requirement that human authorship must occur after AI generated output has been produced in order for a work to be copyrightable.

---

mere mechanical process that operates randomly or automatically without any ***creative input or intervention from a human author")(*** emphasis added).

[4] See, e.g, 2023 Copyright Office Guidance. Part III ("in the case of works containing AI-generated material, the Office will consider whether the AI contributions are the result of 'mechanical reproduction' or instead of an author's 'own original mental conception, to which [the author] gave visible form.' The answer will depend on the circumstances, particularly how the AI tool operates and how it was used to create the final work. This is necessarily a case-by-case inquiry.")

Infringement

We believe in respecting the rights of authors of the underlying data used to train AI models. If a model's output is substantially similar to copyrighted training data, that should be considered infringement, with potential liability for whoever fixes the output in a tangible medium (whether the model creator or the end user) without authorization. This would incentivize AI model creators to obtain the rights to their underlying training data when necessary, and to build models that are less likely to output their training data verbatim. It would also provide for infringement when an AI model end user engineers a prompt that results in infringing output.

Thank you for the opportunity to submit these comments. We welcome further dialogue with the Copyright Office on these important issues.

Sincerely,

rar---

Jennifer Pascua
Vice President Global IP
Jpascua@illumina.com

Illumina • 5200 Illumina Way • San Diego, CA 92122 • Tel 858.202.4500 • Fax 858.202.4766 • www.illumina.com

Comment from AEM Services D/B/A Joseph Street Digest

Comment ID: COLC-2023-0006-7304

I'm answering question #5 concerning proposals for legislative items.
I believe there is new legislation that is needed for the U.S. to get a head of the situation that AI presents regarding copyrights.
I'm proposing the following legislative ideas:
1.AI generated materials define any output from a computer program where a human entered prompt is used, and the AI uses a database or large language system it's been trained on to accomplish this generation output.
2.AI assisted materials are items in which human interaction is still most of the input, but AI acts only in a limited sense such providing suggestive sentence improvements or limited artistic transformations in the part of a digital image.
3.Produced by AI would mean it has altered digital items 60% or more of the digital work in question. Human authors will still produce anything less than this.
4.AI produced works, whether AI generated or AI assisted, maybe copyright protected with an AI Voucher Registration.
5.AI Voucher Registration is only valid for 25 years from publication. Both registration and publication are required for protection. Automatic protections are not granted to works simply printed or published on the internet. Retroactive registration will be granted upon application of works created before legislative adoption.
6.All AI generative software or assistive technology may not enforce a copyright upon its licensed users, regardless of fees charged or not.
7.All AI generative software or assistive technology must get the rights to any works, and/or systems they have used to allow users to produce their works. This will include works in the public domain if such works still have entities maintaining some control over them.
8.#7 would apply for private systems developed and used by any businesses seeking to take advantage of AI generative software or assistive technology for profit.
9.Human AI Endowment Trust payment will be collected by the United States Copyright Office, with each registration equal to $10 USD. The excise tax funds scholarships for students with artistic or editorial abilities.

Comment from Getty Images (US) Inc.

Comment ID: COLC-2023-0006-10260

195 Broadway, 10th Floor, New York, NY 10007
(800) IMAGERY (800) 462-4379 | www.gettyimages.com

legal@gettyimages.com for related inquiries

December 5, 2023

Reply Comments for USCO Inquiry on Artificial Intelligence and Copyright

Docket No. 2023-6

Getty Images appreciates the opportunity to respond to the initial round of comments submitted in response to the U.S. Copyright Office's Notice of Inquiry on Artificial Intelligence and Copyright (the "NOI"). As with our initial response to the NOI ("Initial Comments"), our reply comments focus on Generative AI Models and Generative AI Systems that are trained with copyrighted images and video and the corresponding captions and metadata. In particular, we focus on arguments offered in support of the claims that:

(i) developers of Generative AI Models and Generative AI Systems do not infringe copyrights when they reproduce and distribute copyrighted materials without permission from the copyright owners in the course of training models, even when those models and systems are deployed for commercial purposes and generate outputs that compete with the works on which they are trained, and

(ii) when Generative AI Systems deliver outputs that are substantially similar to copyrighted works on which the underlying model was trained, and therefore infringing, the users are solely to blame and there should be no liability for the model developers.

Both of these propositions are incorrect as a matter of copyright law and misguided as a matter of copyright policy.

That the copyrighted materials on which a Generative AI Model was trained were "publicly available" is not a defense to copyright infringement.

A number of responders observed in their initial comments that their Generative AI Models were trained on "publicly available information." But being "publicly available" and being in the "public domain" are very different propositions. The public availability of a work is irrelevant to the questions of whether the work is copyrighted and whether reproducing the work in the course of training a Generative AI Model without the permission of the copyright owner constitutes copyright infringement. Getty Images does not waive its copyrights or grant any implied licenses by making its copyrighted images available (albeit with prominent watermarks and copyright notices to deter infringement) for review by customers and prospective customers on its public-facing websites. Nor does Getty Images waive its copyrights when it grants a license to a customer to display one of its images on the customer's public-facing website. While the unauthorized reproduction of copyrighted works from a copyright owner's own website or another source authorized by the copyright owner to display the works infringes the copyright owner's exclusive rights absent a fair use defense or other legal justification, we note that some commenters are not only conflating "publicly available" with "public domain" but going so far as to include within the scope of "publicly available information" both works that are obtained from known piracy sites and works scraped in violation of express terms of use. While the public may be able to access such works, they are not "publicly available" for the taking. Courts routinely find defendants liable for infringement when they copy "publicly available" copyrighted photographs from websites for commercial purposes without permission from the copyright owner.[1]

The Copyright Office should emphatically reject efforts to elide the distinction between works that are "publicly available" and works that are in the "public domain" for purposes of assessing whether unauthorized copying is legally defensible.

Training a Generative AI Model – at least in the context of text-to-image or image-to-image generation – involves reproducing the expressive content of the materials on which the model is trained.

Some commenters have attempted to claim that training a Generative AI Model does not even involve acts of copying within the meaning of the Copyright Act. While there may be room to debate, in some circumstances, whether particular acts of copying can fall within the scope of a valid fair use defense, there is no question that training a Generative AI Model capable of text-to-image or image-to-image generation involves reproducing the expressive content of the images and associated captions and metadata on which the model is trained. Indeed, training a diffusion model capable of image generation typically requires at least the following acts of reproduction:

- Copying text-and-image pairings that the developer selects for its training set from the sources from which those pairings are obtained;
- Loading those text-and-image pairings into computer memory to train the model;
- Encoding the images to create smaller versions of the images that take up less memory and separately encoding the paired text, with the encoded images and text retained and stored as an essential element of training; and

[1] *See, e.g., Mango v. Buzzfeed*, *Inc.*, 970 F.3d 167 (2d Cir. 2020); *FameFlynet, Inc. v. Shoshana Collection, LLC*, 282 F. Supp. 3d 618 (S.D.N.Y. 2017); *BWP Media USA Inc. v. HipHopzilla, Inc.*, No. 1:14-CV-0016-AT, 2016 WL 4059682 (N.D. Ga. Mar. 1, 2016).

- Incrementally adding visual “noise” to the encoded images so that it is incrementally harder to discern what is represented and, in the process, creating new copies that, notwithstanding the addition of noise, are substantially similar to the original images (at least until so much noise has been added that they are no longer recognizably similar).

In each instance, the model developer is copying expressive content, not merely some fact or idea contained in the original work.

The expressive components of Getty Images’ copyrighted works are why those works are particularly desirable for use in connection with training Generative AI Models capable of image generation: the photographs are high quality, and the accompanying captions and metadata are richly descriptive. For example, if one were to use the text-image pairing set forth below to train a Generative AI Model, the photograph is much more helpful to the training process than a stick-figure drawing of a dog or a woman with a dog. And it is much more helpful to inform the model that this is a “[p]ortrait of a beautiful young Asian woman walking with her dog at the park under beautiful sunshine” and that this image is a reflection of someone “[s]haring a life with a pet” than merely that this is a picture of a “dog” or a “woman and a dog.”

Similarly, if one were to use the text-image pairing set forth below to train a Generative AI Model, the photograph is much more helpful to the training process than a stick figure of a woman and a young girl or a photograph of a woman and a young girl that does not have the rich context of the surroundings and background. And it is much more helpful to inform the model that this is an image of a "[j]oyful young Asian mother and lovely daughter traveling by airplane," that the "little girl is rolling a suitcase at airport terminal while waiting for departure," and that it reflects both "[f]amily travel and vacation" and "[e]mbarking on a new journey with the family," than merely that this is a picture of a "woman and her daughter" or a picture of "a girl rolling a suitcase."

Copying the expressive content of the text-image pairing is a critical driver of the effectiveness of the training. Moreover, it is not just the expressive content of the image and the expressive content of the text, but how those two forms of expression relate to one another. Consider the following text-image pairing for use in training a Generative AI Model:

Informing the model that "Russian Prime Minister Vladimir Putin is pictured with a horse during his vacation outside the town of Kyzyl in Southern Siberia on August 3, 2009" is far more useful than describing this image as "a man with a horse" or "a picture of Vladimir Putin."

Any contention that training a Generative AI Model does not involve copying of expressive content is simply false.

Sony Corporation of America v. Universal City Studios, Inc. and the "staple article of commerce" doctrine do not excuse infringement arising out of the development and use of Generative AI Models and Generative AI Systems.

A number of commenters contend that Generative AI Models and Generative AI Systems are "staple articles of commerce" and therefore model developers are absolved from liability in light of the U.S. Supreme Court's decision in *Sony Corporation of America v. Universal City Studios, Inc.*, 464 U.S. 417 (1984). These contentions ignore fundamental differences between the nature of copyright infringement claims copyright owners have asserted against developers of Generative AI Models and Generative AI Systems on the one hand and claims asserted against manufacturers of tools considered to be "staple articles of commerce" on the other.

Unlike the training of Generative AI Models and the offering of Generative AI Systems, the manufacture and distribution of the Betamax video recorder at issue in *Sony* did not involve *any* copying of copyrighted works. The only copying at issue occurred *after* the recorders had been sold to consumers, and Sony exercised no control over how consumers chose to use the recorders they owned. The plaintiffs could not assert direct infringement claims against Sony because Sony had not engaged in any copying; they could only assert claims for secondary liability to try to hold Sony accountable for allegedly infringing activities by Betamax users. Adopting the "staple article of commerce" from patent law, the Court held that the Betamax video recorder could be used for "commercially significant non-infringing purposes" and, therefore, Sony was not liable for secondary infringement on the facts presented.[2] The same is true for a variety of other technological tools that can be used to facilitate infringement, but are not themselves infringing, such as the typewriter, the camera, the photocopier, and the printer. In each instance, the manufacturer does not violate any of the exclusive rights afforded to copyright owners in the course of developing or distributing the product, and it does not exercise any control over what consumers do with the product after it has been sold.

Developers of Generative AI Models and Generative AI Systems, in marked contrast, do tread on exclusive rights afforded to copyright owners under Section 106 of the Copyright Act both when they reproduce those works in the course of training models and if they generate and distribute infringing output. That is why, in the copyright infringement cases brought to date against model developers, plaintiffs have sued developers for *direct* infringement. Some plaintiffs may *also* have sued under theories of secondary liability for infringement by users or commercial partners, but the "staple article of commerce" doctrine does not provide a defense to acts of direct infringement. Nor does the doctrine provide a defense to claims of vicarious infringement.[3] The "staple article of commerce" defense only applies to claims of contributory infringement and, even there, model developers are unlikely to qualify for such a defense given the nature of their offerings and the control they maintain when distributing those offerings.

In most instances to date, model developers are not simply injecting their products into the stream of commerce with no knowledge of how their products are used and no ability to

---

[2] *Sony*, 464 U.S. at 442.

[3] *See, e.g.*, *A&M Records, Inc. v. Napster, Inc.*, 239 F.3d 1004, 1022-23 (9th Cir. 2001) ("*Sony*'s 'staple article of commerce' analysis has no application to Napster's potential liability for vicarious copyright infringement.").

control that use. To the contrary, most developers of Generative AI Systems are selling consumers and businesses *access* to the use of an underlying model, typically in the form of a subscription to one of the services they offer.[4] And this is hardly a passive engagement. Numerous comments submitted on behalf of developers of Generative AI Models and Generative AI Systems touted the steps those entities are taking to ensure the safety of the products and the guardrails they have in place, whether in relation to mitigating the likelihood of delivering infringing outputs or, more commonly, a variety of other harms unrelated to copyright. Numerous courts have found the distinction between selling a product and offering an ongoing service critical for purposes of evaluating whether the "staple article of commerce" doctrine provides a defense to a claim of contributory infringement.[5]

Accordingly, to the extent that some commenters have suggested that the "staple article of commerce" doctrine absolves model developers from potential liability for copyright infringement, they are mistaken.

By selecting copyrighted works for inclusion in training sets, and reproducing such works in the course of training, model developers have engaged in "volitional conduct" for purposes of copyright law.

Some commenters have attempted to suggest that entities that develop and offer Generative AI Systems cannot be liable for copyright infringement because they purportedly do not engage in "volitional conduct." These arguments are specious.

As an initial matter, these arguments do not apply to infringement claims based on unauthorized copies made in the course of training a Generative AI model. Such copying is deliberate and intentional. Developers choose what works they will copy and the sources from

---

[4] Such subscriptions are typically governed by contractual agreements between the user and the distributor of the service (whether in the form of a "click through" agreement or otherwise) that impose limitations on what the user can do with the service and reserve important rights for the distributor.

[5] *See, e.g.*, *Arista Records LLC v. Usenet.com, Inc.*, 633 F. Supp. 2d 124, 156 (S.D.N.Y.2009) (rejecting defendants invocation of *Sony* as a complete defense and finding non-infringing uses "immaterial" where there was an ongoing relationship with users); *see also Capitol Records, LLC v. ReDigi Inc*, 934 F. Supp. 2d 640, 658-59 (S.D.N.Y. 2013) ("*ReDigi*"), *aff'd*, 910 F.3d 649 (2d Cir. 2018) (finding, in addition to liability for direct infringement for defendant's own acts of reproduction and distribution, liability for contributory infringement where service "provided the site and facilities for their users' infringement") (cleaned up).

which they will obtain those copies. While a plaintiff may need discovery to know specifically which entity or individual was responsible for those choices, there is no room to dispute that whoever was engaged in the actual acts of reproduction did so with requisite volitional conduct for infringement liability, regardless of whether they understood their acts to be unlawful at the time. Copyright infringement is a strict liability tort.

There is ample volitional conduct for liability arising out of claims based on the delivery of infringing outputs as well, even though the user, rather than the developer, selects the prompts. But the prompts alone do not determine the output, as evidenced by the fact that a Generative AI System will typically deliver a variety of different outputs in response to the same set of prompts. And even if they did, in each instance, the output will reflect of how the underlying model was trained, including the choices that the developer made with respect to what works to include the training set, where and how to obtain copies of those works, what type of fine-tuning to provide, the extent of fine-tuning, whether to implement filters of datasets and, if so, which filters to use.[6] The decision to use copyrighted works to train a Generative AI Model and the selection of which works to use distinguish such models from other technological tools in which volitional conduct by the offeror was found lacking, even though users could use those tools to infringe.[7] Model developers are not "passive conduits"; their active engagement in the process by which a user is supplied with an infringing output supplies ample volitional conduct.[8] By selecting particular prompts, users also play a role that influences what content is delivered to them, but that does not absolve offerors of Generative AI Systems from liability for making and delivering an infringing output.[9]

---

[6] Several commenters (which are also defendants in pending copyright infringement lawsuits) touted that they filter datasets for issues like "safety," "bias," and "quality" or "conduct legal and ethical diligence on the data that we use." The decision not to filter datasets to exclude unlicensed copyrighted works is deliberate.

[7] *Compare Cartoon Network, LP v. CSC Holdings, Inc.*, 536 F.3d 121, 131 (2d Cir. 2008) ("*Cablevision*") *with Fox News Network, LLC v. Tveyes, Inc.*, 883 F.3d 169, 181 (2d Cir. 2018) (distinguishing *Cablevision*) and *ReDigi*, 934 F. Supp. 2d at 656-57 (same).

[8] *See Arista Records*, 633 F. Supp. 2d at 148-49.

[9] *See Am. Broad Co., Inc. v. Aereo, Inc.*, 573 U.S. 431 (2014) (reversing appellate court's determination that provider of service that facilitated subscriber access to Internet streams of copyrighted broadcast

The number of works that a developer may wish to use to train a Generative AI Model does not excuse unauthorized copying.

A number of commenters have complained that licensing the use of copyrighted works in training sets would be either impossible, impractical, or unduly expensive because of the sheer number of works some model developers would like use for training purposes. The scope of infringement in which an infringer would like to engage hardly excuses the infringement.

Voluntarily offering "opt outs" from the unauthorized use of copyrighted materials in training sets does not excuse infringement.

A number of responders who have engaged in the unauthorized use of copyrighted materials to train Generative AI Models and Generative AI Systems recognize that, generally speaking, the creative community vehemently objects to these uses and now claim to want to be sensitive to creator concerns. They suggest that a few half-measures, such as permitting copyright owners to "opt out" of their unauthorized uses, should be sufficient to resolve any objections to their infringing conduct.

Getty Images addressed the numerous reasons why an "opt-out" approach to AI training is not a suitable substitute for affirmative consent in our Initial Comments (at 18-19). But, in any event, forward-looking opt-out regimes are "too little too late" to excuse the infringement to date. While Getty Images of course welcomes reasonable dialogue with model developers on these issues, it is not credible for developers to feign surprise that copyright owners would object to the commercial exploitation of their copyrighted works in this manner, and so even if an opt-out scheme could have passed muster as a matter of copyright law—and it cannot—it should have been offered from the start. Moreover, all of these programs are (at best) voluntary. Other than copyright law, there is nothing stopping any of the commenters who have described such offers from changing their policies tomorrow and refusing to honor "opt out" requests.[10] Toothless commitments such as these should have no bearing on how the

---

television programming had not engaged in "volitional conduct" and holding provider liable for direct infringement).

[10] It is not even clear whether such offers will be rigorously and scrupulously applied. The Copyright Office, at a minimum, should consider whether and to what extent some model developers might honor requests from copyright owners whose works they do not view as important to the training or for which

Copyright Office, the courts, and Congress consider the interplay of artificial intelligence and copyright law.

These voluntary "opt-out" programs are cold comfort to copyright owners such as Getty Images whose works *already* have been used extensively without its permission. There have been no offers to cease distributing models that were trained on our works without authorization, no offers to unwind the effects of training on such works, and no explanation of how it would even be possible to do so. Absent commitments to train future models and future versions on entirely new sets of inputs that do not include "opt out" works, rather than building on already released models to develop future ones, these "opt out" offers are even less meaningful.

Robots.txt is not an effective way for copyright owners to prevent the unauthorized use of their copyrighted materials to train Generative AI Models.

A number of commenters have suggested that there is little reason to be concerned with unauthorized use of copyrighted materials in training sets for Generative AI Models because copyright owners could indicate a "do not train" preference in their website robot.txt settings. In its Initial Comments (at 19), Getty Images identified several reasons why reliance on such indications of preferences is inadequate, including that (i) it may require copyright owners to block web crawling that they do find desirable, such as (in our case) the indexing of websites by search engines to make them easier for users to find, and (ii) many copyright owners, including Getty Images, license the display of their copyrighted works on third-party websites and do not control what preferences those licensees indicate with respect to web crawling on their own websites.[11] In addition, such instructions would do nothing to address the gobsmacking extent of unauthorized copying that has already taken place.

there are easy substitutes, but then refuse to honor requests from large scale copyright owners with particularly important collections of works, such as Getty Images.

[11] Nor could Getty Images and other copyright licensors reasonably be expected to dictate website settings to their customers as a condition of licensing, especially since those settings are set at a website level, rather than applied to individual pieces of content. In any event, there is no basis to impose the incredibly disruptive burden of trying to negotiate, police, and enforce such conditions on copyright owners in order to protect the exclusive rights copyright law provides, and there is no reason to stifle the licensing that would occur but for the imposition of such a requirement in the event customers are unwilling to accept conditions of that nature.

Other comments submitted in response to the NOI betray an even more fundamental problem with relying on robots.txt instructions as a solution to the intrusion on copyright owners' exclusive rights: those instructions can be ignored. Other comments admit that some model developers are willing to violate website terms of use that expressly prohibit copying for commercial purposes and—worse—are willing to obtain copies of works from known piracy sites to facilitate training Generative AI Models because they think the fair use doctrine will still countenance their unauthorized copying in this context.[12] Because compliance with the Robots Exclusion Protocol is voluntary, a robots.txt instruction would be, at most, a minor speedbump for such unscrupulous actors, no different than the other unmistakably clear indications of copyright owner preferences that they have already willfully ignored. For example, Getty Images already (i) expressly prohibits visitors to its websites from downloading, copying, or re-transmitting any of its content and from using data mining, robots or similar data gathering or extraction methods, (ii) places conspicuous watermarks on each of its images and videos displayed on those websites in order to deter infringement, and (iii) displays conspicuous copyright notices on its websites. Is anyone naïve enough to think that Stability AI would have been deterred by a robots.txt instruction when it decided to scrape 12 million copyrighted images from Getty Images websites from its pre-determined list without permission as part of developing its Stable Diffusion Generative AI Model when it was not deterred by any of the other measures Getty Images had already taken to prevent unauthorized copying?

Developers of Generative AI Models and Generative AI Systems cannot absolve themselves of liability for unauthorized copying with facile but inapt analogies to "human learning."

A number of commenters observe that human beings learn by reading, listening to, or looking at copyrighted materials and often build on what they have learned by creating new works that do not infringe on the source materials from which they learned. In their view, the unauthorized acts of copying and distribution associated with Generative AI Models and Generative AI Systems can be excused with facile analogies to human learning. These analogies

---

[12] For example, as reported in the *Washington Post*, one data set (Google's C4 data set) used to train prominent large language models included materials obtained from at least 28 different web sites identified by the Office of the U.S. Trade Representative as notorious markets for counterfeiting and piracy. Kevin Schaul, Szu Yu Chen & Nitasha Tiku, *Inside the Secret List of Websites That Make AI Like ChatGPT Sound Smart*, WASH. POST, April 19, 2023 (available at https://www.washingtonpost.com/technology/interactive/2023/ai-chatbot-learning/).

ignore fundamental differences between training a Generative AI Model and "human learning" that are highly consequential for copyright law.

At the risk of belaboring the obvious, Section 106 provides copyright owners with specific exclusive rights, including, for example, exclusive rights of reproduction, distribution, performance, and display.[13] Copyright law does not provide copyright owners with exclusive rights of learning, reading, looking at, or listening to. While using copyrighted works to train a Generative AI Model involves numerous acts within the scope of Section 106[14], the "human learning" to which these commenters analogize does not implicate any of the exclusive rights belonging to the copyright owner. When humans do tread on a copyright owner's exclusive right in the course of learning or research, their otherwise infringing acts will sometimes be defensible as a fair use.[15] But they are not necessarily so. Indeed, the scientific publishing and educational publishing industries (among many others) depend on the notion that copyrighted materials cannot be indiscriminately copied and shared without permission from the copyright owner, just because the copying is intended to facilitate learning or research.[16]

Developers of Generative AI Models should be required to prepare, retain, and disclose auditable records of the materials used to train their models and the sources from which those materials were obtained.

In its Initial Comments (at 23-24), Getty Images explained why developers of Generative AI Models should be required to collect, retain, and disclose records regarding the materials used to train their models on the grounds that the training sets that they have elected to use are proprietary. Getty Images also described (at 23-24) the level of specificity that we believe should be required.

---

[13] 17 U.S.C. § 106.

[14] *see* pp. 3-9 *supra* and Initial Comments at 9-11

[15] 17 U.S.C. §107.

[16] *See, e.g.*, *Princeton Univ. Press v. Michigan Document Servs., Inc.*, 99 F.3d 1381 (6th Cir. 1996); *American Geophysical Union v. Texaco*, *Inc.*, 60 F.3d 913 (2d Cir. 1995); *Basic Books, Inc. v. Kinko's Graphics Corp.*, 758 F. Supp. 1522 (S.D.N.Y. 1991).

Some entities that offer Generative AI Models and Generative AI Systems agree that at least some form of transparent, public disclosure of training data should be required.[17] Others, however, oppose required disclosures on the ground that training sets are trade secrets or are otherwise proprietary. The Copyright Office should reject such efforts to exalt whatever purported intellectual property interests model developers may have in the catalog of copyrighted works they have purloined to train their models over the copyright interests in the works on which models were trained. Some commenters have also noted that a disclosure required could tread on confidentiality or privacy interests. Getty Images believes that a disclosure requirement could interfere with *bona fide* privacy interests in only the most exceptional cases. Such cases would be better handled through authorization to redact truly private information (such as social security numbers, checking account numbers, personally identifiable information contained in medical records, and the like) from the required disclosures, rather than by dispensing with a disclosure requirement entirely.

****

We are thankful to the USCO for providing the opportunity to respond to the initial round of comments submitted in response to the NOI. As expressed in our Initial Comments, while we do not think it is necessary to amend existing copyright law to address issues raised by AI, we do believe that transparency standards are necessary to protect copyright interests. We are highly supportive of regulatory efforts to mandate a set of consistent standards applicable to the development and deployment of AI and believe that such standards are key to promoting responsible innovation and global harmonization. We look forward to being part of the solution to ensure that the AI ecosystem prospers, while respecting existing intellectual property rights.

Respectfully submitted,

Getty Images

---

[17] *See, e.g.*, Hugging Face Response to the Copyright Office Notice of Inquiry on Artificial Intelligence and Copyright at 2, 4, 12.